

modelling_naive

September 23, 2019

```
[1]: import pickle

%matplotlib inline
import matplotlib.pyplot as plt

import numpy as np
import pandas as pd
pd.options.mode.chained_assignment = None

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import StratifiedKFold, StratifiedShuffleSplit, \
    train_test_split, learning_curve
from sklearn.metrics import classification_report, confusion_matrix

from sklearn.linear_model import LogisticRegression

from lightgbm import LGBMClassifier

[2]: # Load preprocessed dataset

PREPROCESSED_DF_PATH = 'preprocessed_applications_df.dataframe.pd'
main_df = pickle.load( open(PREPROCESSED_DF_PATH, 'rb') )

[3]: # Leave only 'naive' features: age, gender, region, monthlyincome, monthlycost, \
    education, family status

naive_features = [
    'days_from_birth',
    'gender',
    'region', 'binned_region',
    'monthlyincome', 'monthlycost',
    'fam_status', 'binned_fam_status',
    'df' # target featur
]
naive_df = main_df[ naive_features ]
```

```
[4]: # Convert several columns to 'category' dtype

category_features = ['gender', 'binned_region', 'binned_fam_status']

for feature_name in category_features:
    naive_df[feature_name] = naive_df[feature_name].astype('category')

[5]: # Apply standard scaler (z-transform)

standard_scl = StandardScaler()

features_to_scale = ['days_from_birth', 'monthlyincome', 'monthlycost']
scaled_data = standard_scl.fit_transform( naive_df[features_to_scale] )

naive_df.loc[:, 'days_from_birth'] = scaled_data[:, 0]
naive_df.loc[:, 'monthlyincome'] = scaled_data[:, 1]
naive_df.loc[:, 'monthlycost'] = scaled_data[:, 2]

[6]: # Encode binned features

features_to_encode = ['gender', 'binned_region', 'binned_fam_status']
naive_df = pd.get_dummies( naive_df, columns=features_to_encode,
    ↳drop_first=True )

# Encode train label
train_labels = naive_df['df'].map({
    'bad': 0, 'good': 1
}).dropna()

[7]: def draw_learning_curve(estimator, X_tr, y_tr):
    train_sizes, train_scores, val_scores = learning_curve(
        estimator, X_tr, y_tr, train_sizes=np.linspace(0.1, 1.0, 4), cv=3
    )
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
    val_scores_mean = np.mean(val_scores, axis=1)
    val_scores_std = np.std(val_scores, axis=1)
    plt.grid()
    plt.fill_between(
        train_sizes,
        train_scores_mean - train_scores_std,
        train_scores_mean + train_scores_std,
        alpha=0.1, color="r"
    )
    plt.fill_between(
        train_sizes,
        val_scores_mean - val_scores_std,
        val_scores_mean + val_scores_std,
```

```

        alpha=0.1, color="g"
    )
    plt.plot(
        train_sizes,
        train_scores_mean,
        'o-', color="r", label="Training score"
    )
    plt.plot(
        train_sizes,
        val_scores_mean,
        'o-', color="g", label="Cross-validation score"
    )
    plt.legend(loc="best")
    plt.show()

```

[9]: *# Stratified CV*

```

train_df = naive_df[ naive_df['df'].isnull() == False ]
train_df = train_df.drop( ['df'], axis=1 )

skf = StratifiedKFold(
    n_splits=3,
    shuffle=True
)

for train_idx, val_idx in skf.split(train_df, train_labels):
    X_tr, X_val = train_df.iloc[train_idx, :], train_df.iloc[val_idx, :]
    y_tr, y_val = train_labels.iloc[train_idx], train_labels.iloc[val_idx]

    # Not enough variance to explain the data
    lr_model = LogisticRegression(
        C=0.001,
        solver='lbfgs',
        n_jobs=-1
    )
    lr_model.fit(X_tr, y_tr)
    lr_val_pred = lr_model.predict(X_val)
    print(confusion_matrix(y_val, lr_val_pred))
    draw_learning_curve( lr_model, X_tr, y_tr )

# f1: 0.0:0.2, 1.0:0.8
# lgbm_model = LGBMClassifier(
#     learning_rate=0.05,
#     metric='recall',
#     objective='binary',
#     num_leaves=500,
#     max_depth=250,

```

```


#         n_estimators=1000,
#         scale_pos_weight=5,
#         n_jobs=4,
#     )
#     lgbm_model.fit(X_tr, y_tr)
#     lgbm_val_pred = lgbm_model.predict(X_val)
#     print(confusion_matrix(y_val, lgbm_val_pred))
# #     print(classification_report(y_val, lgbm_val_pred))
#     draw_learning_curve( lgbm_model, X_tr, y_tr )

```

```

[[ 0 330]
 [ 0 1002]]


```




```

[[ 0 330]
 [ 0 1002]]

```



```
[[ 0 330]
 [ 0 1002]]
```



[12]: *# Stratified Shuffle CV*


```
sss = StratifiedShuffleSplit(
    n_splits=3,
    test_size=0.4,
)

for train_idx, val_idx in sss.split(train_df, train_labels):
    X_tr, X_val = train_df.iloc[train_idx, :], train_df.iloc[val_idx, :]
    y_tr, y_val = train_labels.iloc[train_idx], train_labels.iloc[val_idx]


    # Not enough variance to explain the data
    model_logreg = LogisticRegression(
        C=0.04,
        solver='lbfgs',
        max_iter=1000,
        n_jobs=-1
    )
    model_logreg.fit( X_tr, y_tr )
    logreg_val_pred = model_logreg.predict( X_val )
    # print(classification_report(y_val, logreg_val_pred))
    print(confusion_matrix(y_val, logreg_val_pred))
    draw_learning_curve( lr_model, X_tr, y_tr )

# f1: 0.0:0.2, 1.0:0.8
# lgbm_model = LGBMClassifier(
#     learning_rate=0.05,
#     metric='recall',
#     objective='binary',
#     num_leaves=500,
#     max_depth=250,
#     n_estimators=1000,
#     scale_pos_weight=5,
#     n_jobs=4,
# )
# lgbm_model.fit(X_tr, y_tr)
# lgbm_val_pred = lgbm_model.predict(X_val)
# # display( pd.Series(lgbm_val_pred).value_counts() )
# print(confusion_matrix(y_val, lgbm_val_pred))
# # print(classification_report(y_val, lgbm_val_pred))
# draw_learning_curve( lgbm_model, X_tr, y_tr )
```

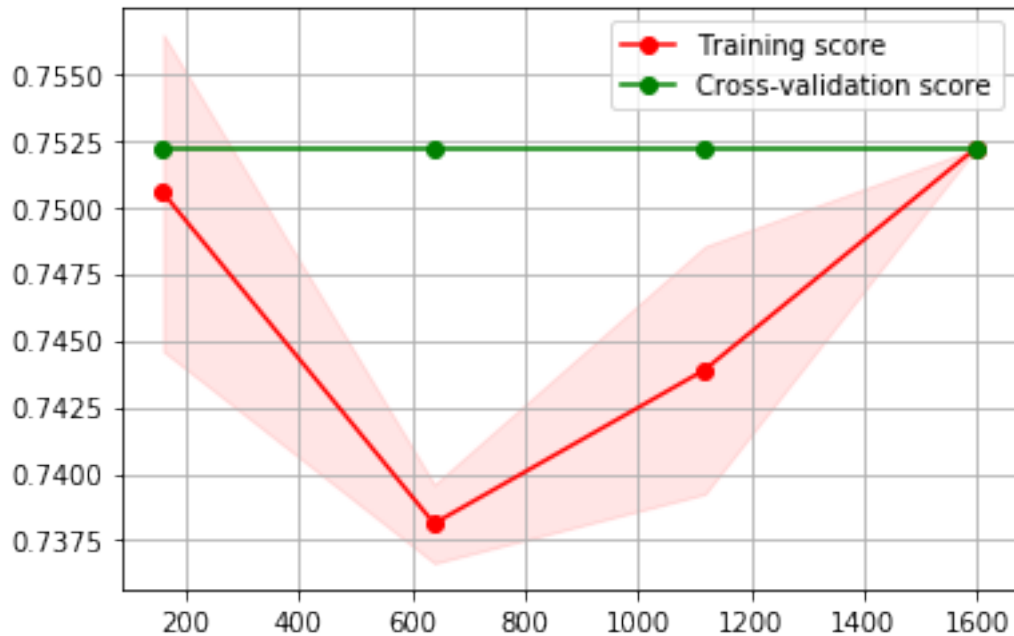
```
[[ 1 395]
 [ 0 1203]]
```



```
[[ 1 395]
 [ 0 1203]]
```



```
[[ 1 395]
 [ 0 1203]]
```



[10]: *# Feature importances*

```
lgbm_model = LGBMClassifier(
    learning_rate=0.05,
    metric='recall',
    objective='binary',
    num_leaves=500,
    max_depth=250,
    n_estimators=1000,
    scale_pos_weight=5,
    n_jobs=4,
)
lgbm_model.fit(X_tr, y_tr)
display(
    pd.DataFrame({
        'imp_name': train_df.columns,
        'imp_value': lgbm_model.booster_.
        →feature_importance(importance_type='gain'),
    }).sort_values('imp_value', ascending=False)
)
```

	imp_name	imp_value
0	days_from_birth	17710.272379
2	monthlyincome	8464.744314
3	monthlycost	7670.087091
1	region	6829.490372

4	fam_status	1571.855154
5	gender_2	1086.629987
8	binmed_fam_status_2	695.506884
7	binmed_region_3	583.874688
6	binmed_region_2	489.663288
9	binmed_fam_status_3	0.000000